


The Data Science Track



Prepared By: R. Daynola

1

1

INTRODUCTION TO DATA SCIENCE WITH R



15.LOOP FUNCTIONS

2

2

LOOPING ON THE COMMAND LINE

Writing for, while loops is useful when programming but not particularly easy when working interactively on the command line. There are some functions which implement looping to make life easier.

- `lapply`: Loop over a list and evaluate a function on each element
- `sapply`: Same as `lapply` but try to simplify the result
- `apply`: Apply a function over the margins of an array
- `tapply`: Apply a function over subsets of a vector
- `mapply`: Multivariate version of `lapply`

An auxiliary function `split` is also useful, particularly in conjunction with `lapply`.



3

3

`lapply`

`lapply` takes three arguments: (1) a list `x`; (2) a function (or the name of a function) `FUN`; (3) other arguments via its `...` argument. If `x` is not a list, it will be coerced to a list using `as.list`.

```
> lapply
function (X, FUN, ...)
{
  FUN <- match.fun(FUN)
  if (!is.vector(X) || is.object(X))
    X <- as.list(X)
  .Internal(lapply(X, FUN))
}
<bytecode: 0x0000000009dfc290>
<environment: namespace:base>
```

Note: The actual looping is done internally in **C** code.



4

4

lapply

`lapply` always returns a list, regardless of the class of the input

```
> set.seed(100)
> x <- list(a = 1:5, b = rnorm(10))
> lapply(x, mean)
$a
[1] 3

$b
[1] -0.01795716
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

5

5

lapply

```
> set.seed(100)
> x <- list(a = 1:4, b = rnorm(10), c = rnorm(20, 1), d = rnorm(100, 5))
> lapply(x, mean)
$a
[1] 2.5

$b
[1] -0.01795716

$c
[1] 1.052275

$d
[1] 4.971873
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

6

6

lapply

```
> set.seed(100)
> x <- 1:4
> lapply(x, runif)
[[1]]
[1] 0.3077661

[[2]]
[1] 0.2576725 0.5523224

[[3]]
[1] 0.05638315 0.46854928 0.48377074

[[4]]
[1] 0.8124026 0.3703205 0.5465586 0.1702621
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

7

7

lapply

```
> set.seed(100)
> x <- 1:4
> lapply(x, runif, min = 0, max = 10)
[[1]]
[1] 3.077661

[[2]]
[1] 2.576725 5.523224

[[3]]
[1] 0.5638315 4.6854928 4.8377074

[[4]]
[1] 8.124026 3.703205 5.465586 1.702621
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

8

8

lapply

`lapply` and friends make heavy use of **anonymous** functions.

```
> x <- list(a = matrix(1:4, 2, 2), b = matrix(1:6, 3, 2))
> x
$a
  [,1] [,2]
[1,]  1   3
[2,]  2   4

$b
  [,1] [,2]
[1,]  1   4
[2,]  2   5
[3,]  3   6
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

9

9

lapply

An anonymous function for extracting the first column of each matrix.

```
> lapply(x, function(elt) elt[,1])
$a
[1] 1 2

$b
[1] 1 2 3
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

10

10

sapply

`sapply` will try to simplify the result of `lapply` if possible.

- If the result is a list where every element is length 1, then a vector is returned
- If the result is a list where every element is a vector of the same length (> 1), a matrix is returned.
- If it can't figure things out, a list is returned



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

11

11

sapply

```
> set.seed(100)
> x <- list(a = 1:4, b = rnorm(10), c = rnorm(20, 1), d = rnorm(100, 5))
> lapply(x, mean)
$a
[1] 2.5

$b
[1] -0.01795716

$c
[1] 1.052275

$d
[1] 4.971873
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

12

12

sapply

```
> sapply(x, mean)
```

	a	b	c	d
	2.50000000	-0.01795716	1.05227455	4.97187322

```
> mean(x)
```

```
[1] NA
```

```
Warning message:
```

```
In mean.default(x) : argument is not numeric or logical: returning NA
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

13

13

apply

`apply` is used to evaluate a function (often an anonymous one) over the margins of an array.

- It is most often used to apply a function to the rows or columns of a matrix
- It can be used with general arrays, e.g. taking the average of an array of matrices
- It is not really faster than writing a loop, but it works in one line!



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

14

14

apply

```
> str(apply)
function (X, MARGIN, FUN, ...)
```

- X is an array
- MARGIN is an integer vector indicating which margins should be “retained”.
- FUN is a function to be applied
- ... is for other arguments to be passed to FUN



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

15

15

apply

```
> set.seed(100)
> x <- matrix(rnorm(200), 20, 10)
> apply(x, 2, mean)
[1] 0.10786715 0.09248710 -0.01819675
[4] -0.10476571 -0.06282897 -0.01056873
[7] -0.02770605 -0.07369189 0.11375866
[10] 0.05391220
> apply(x, 1, sum)
[1] -3.4041391 2.4794830 -4.2674893
[4] 5.1227369 -3.7383995 -1.6227874
[7] -2.6503351 2.4200274 0.0243034
[10] -1.6415171 0.6320444 1.4873183
[13] -1.2966852 1.1269602 -2.3229059
[16] 0.9502779 -0.8370400 3.3275439
[19] 1.4669027 4.1490405
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

16

16

col/row sums and means

For sums and means of matrix dimensions, we have some shortcuts.

- `rowSums = apply(x, 1, sum)`
- `rowMeans = apply(x, 1, mean)`
- `colSums = apply(x, 2, sum)`
- `colMeans = apply(x, 2, mean)`

The shortcut functions are much faster, but you won't notice unless you're using a large matrix.



17

17

OTHER WAYS TO apply

Quantiles of the rows of a matrix.

```
> set.seed(100)
> x <- matrix(rnorm(200), 20, 10)
> apply(x, 1, quantile, probs = c(0.25, 0.75))
```

	[,1]	[,2]	[,3]	[,4]
25%	-0.4861668	-0.02273898	-0.55524707	0.4669936
75%	-0.2472011	0.72051393	0.07328519	0.8308268
	[,5]	[,6]	[,7]	[,8]
25%	-0.9399706	-0.6020301	-0.5456799	-0.3507716
75%	0.1266184	0.1994963	-0.1330212	0.7916547
	[,9]	[,10]	[,11]	[,12]
25%	-0.7363371	-0.2967551	-0.3715975	-0.7004898
75%	0.3648825	0.2188459	0.3798258	1.1348764
	[,13]	[,14]	[,15]	[,16]
25%	-0.5113089	-0.4116403	-0.6500849	-0.3082920
75%	0.1296619	0.6272514	0.2185301	0.3644001
	[,17]	[,18]	[,19]	[,20]
25%	-0.7215854	-0.03367689	-0.7609959	-0.594393
75%	0.3145039	0.44599246	1.0511647	1.453944

ANALYSIS STRUCTURE ALGORITHM PROCESS PROGRAMMING SOLVING KNOWLEDGE

18

18

OTHER WAYS TO apply

Average matrix in an array.

```
> set.seed(100)
> a <- array(rnorm(2 * 2 * 10), c(2, 2, 10))
> apply(a, c(1, 2), mean)
      [,1]      [,2]
[1,] -0.4166189 -0.1535242
[2,]  0.2219812  0.7488705
> rowMeans(a, dims = 2)
      [,1]      [,2]
[1,] -0.4166189 -0.1535242
[2,]  0.2219812  0.7488705
```



19

19

mapply

`mapply` is a multivariate apply of sorts which applies a function in parallel over a set of arguments.

```
> str(mapply)
function (FUN, ..., MoreArgs = NULL, SIMPLIFY = TRUE,
        USE.NAMES = TRUE)
```

- `FUN` is a function to apply
- `...` contains arguments to apply over
- `MoreArgs` is a list of other arguments to `FUN`.
- `SIMPLIFY` indicates whether the result should be simplified



20

20

mapply

The following is tedious to type

```
list(rep(1, 4), rep(2, 3), rep(3, 2), rep(4, 1))
```

Instead we can do

```
> mapply(rep, 1:4, 4:1)
[[1]]
[1] 1 1 1 1

[[2]]
[1] 2 2 2

[[3]]
[1] 3 3

[[4]]
[1] 4
```

ANALYSIS STRUCTURE ALGORITHM PROCESS PROGRAMMING SOLVING KNOWLEDGE

21

21

VECTORIZING A FUNCTION

```
> set.seed(100)
> noise <- function(n, mean, sd){
+   rnorm(n, mean, sd)
+ }
> noise(5, 1, 2)
[1] -0.004384701 1.263062331 0.842165820
[4] 2.773569619 1.233942541
> noise(1:5, 1:5, 2)
[1] 1.6372602 0.8364186 4.4290654 2.3494811
[5] 4.2802757
```



22

22

INSTANT VECTORIZATION

```
> set.seed(100)
> mapply(noise, 1:5, 1:5, 2)
[[1]]
[1] -0.004384701

[[2]]
[1] 2.263062 1.842166

[[3]]
[1] 4.773570 3.233943 3.637260

[[4]]
[1] 2.836419 5.429065 2.349481 3.280276

[[5]]
[1] 5.179772 5.192549 4.596732 6.479681 5.246759
```



23

23

mapply

Which is the same as:

```
> set.seed(100)
> list(noise(1, 1, 2), noise(2, 2, 2),
+      noise(3, 3, 2), noise(4, 4, 2),
+      noise(5, 5, 2))
[[1]]
[1] -0.004384701

[[2]]
[1] 2.263062 1.842166

[[3]]
[1] 4.773570 3.233943 3.637260

[[4]]
[1] 2.836419 5.429065 2.349481 3.280276

[[5]]
[1] 5.179772 5.192549 4.596732 6.479681 5.246759
```



24

24

tapply

tapply is used to apply a function over subsets of a vector.

```
> str(tapply)
function (X, INDEX, FUN = NULL, ..., default = NA,
         simplify = TRUE)
```

- X is a vector
- INDEX is a factor or a list of factors (or else they are coerced to factors)
- FUN is a function to be applied
- ... contains other arguments to be passed FUN
- simplify, should we simplify the result?



25

25

tapply

Take group means.

```
> x <- c(rnorm(10), runif(10), rnorm(10, 1))
> f <- gl(3, 10)
> f
[1] 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 3 3 3 3 3 3 3 3 3 3
Levels: 1 2 3
> tapply(x, f, mean)
      1      2      3
0.006836221 0.459535468 1.395102751
```



26

26

tapply

Take group means without simplification.

```
> tapply(x, f, mean, simplify = FALSE)
$`1`
[1] 0.006836221

$`2`
[1] 0.4595355

$`3`
[1] 1.395103
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

27

27

tapply

Find group ranges.

```
> tapply(x, f, range)
$`1`
[1] -1.878656  1.403203

$`2`
[1] 0.0115455 0.9711167

$`3`
[1] -0.3988256  3.5819589
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

28

28

split

`split` takes a vector or other objects and splits it into groups determined by a factor or list of factors.

```
> str(split)
function (x, f, drop = FALSE, ...)
```

- `x` is a vector (or list) or data frame
- `f` is a factor (or coerced to one) or a list of factors
- `drop` indicates whether empty factors levels should be dropped



29

29

split

```
> x <- c(rnorm(10), runif(10), rnorm(10, 1))
> f <- gl(3, 10)
> split(x, f)
$`1`
[1] -0.62979029 -0.25248978 -0.69042216  0.20254215  0.84638144
[6]  0.63207406  0.20141352 -0.09107064  0.28948413 -0.05468494

$`2`
[1] 0.02058322 0.17096903 0.63996650 0.16497848 0.35472278
[6] 0.18642647 0.89765616 0.23778386 0.98494348 0.02133074

$`3`
[1] -0.23972284 1.58987389 1.12401929 0.47629221 1.62022800
[6] 1.70822158 0.90680165 0.70480330 -0.08581523 0.37518494
```



30

30

split

A common idiom is `split` followed by `lapply`.

```
> lapply(split(x, f), mean)
$`1`
[1] 0.04534375

$`2`
[1] 0.3679361

$`3`
[1] 0.8179887
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

31

31

SPLITTING A DATA FRAME

```
> head(airquality)
  Ozone solar.R Wind Temp Month Day
1   41    190  7.4   67     5   1
2   36    118  8.0   72     5   2
3   12    149 12.6   74     5   3
4   18    313 11.5   62     5   4
5   NA     NA 14.3   56     5   5
6   28     NA 14.9   66     5   6
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

32

32

SPLITTING A DATA FRAME

```
> s <- split(airquality, airquality$Month)
> lapply(s, function(x) colMeans(x[, c("Ozone", "solar.R", "wind")]))
```

\$ 5`

Ozone	solar.R	Wind
NA	NA	11.62258

\$ 6`

Ozone	solar.R	Wind
NA	190.16667	10.26667

\$ 7`

Ozone	solar.R	Wind
NA	216.483871	8.941935

\$ 8`

Ozone	solar.R	Wind
NA	NA	8.793548

\$ 9`

Ozone	solar.R	Wind
NA	167.4333	10.1800

ANALYSIS STRUCTURE ALGORITHM PROCESS PROGRAMMING SOLVING KNOWLEDGE

33

SPLITTING A DATA FRAME

```
> sapply(s, function(x) colMeans(x[, c("Ozone", "solar.R", "wind")]))
```

	5	6	7	8	9
Ozone	NA	NA	NA	NA	NA
solar.R	NA	190.16667	216.483871	NA	167.4333
Wind	11.62258	10.26667	8.941935	8.793548	10.1800

```
> sapply(s, function(x) colMeans(x[, c("Ozone", "solar.R", "wind")], na.rm = TRUE))
```

	5	6	7	8	9
Ozone	23.61538	29.44444	59.115385	59.961538	31.44828
solar.R	181.29630	190.16667	216.483871	171.857143	167.43333
Wind	11.62258	10.26667	8.941935	8.793548	10.18000

ANALYSIS STRUCTURE ALGORITHM PROCESS PROGRAMMING SOLVING KNOWLEDGE

34

SPLITTING MORE THAN ONE LEVEL

```
> x <- rnorm(10)
> f1 <- gl(2, 5)
> f2 <- gl(5, 2)
> f1
[1] 1 1 1 1 1 2 2 2 2 2
Levels: 1 2
> f2
[1] 1 1 2 2 3 3 4 4 5 5
Levels: 1 2 3 4 5
> interaction(f1, f2)
[1] 1.1 1.1 1.2 1.2 1.3 2.3 2.4 2.4 2.5 2.5
Levels: 1.1 2.1 1.2 2.2 1.3 2.3 1.4 2.4 1.5 2.5
```



35

35

SPLITTING MORE THAN ONE LEVEL

Interactions can create empty levels.

```
> str(split(x, list(f1, f2)))
List of 10
 $ 1.1: num [1:2] -0.233 -0.251
 $ 2.1: num(0)
 $ 1.2: num [1:2] 0.954 -0.266
 $ 2.2: num(0)
 $ 1.3: num 1.9
 $ 2.3: num -0.43
 $ 1.4: num(0)
 $ 2.4: num [1:2] 1.576 0.162
 $ 1.5: num(0)
 $ 2.5: num [1:2] -1.085 0.577
```



36

36

SPLITTING MORE THAN ONE LEVEL

Empty levels can be dropped.

```
> str(split(x, list(f1, f2), drop = TRUE))
List of 6
 $ 1.1: num [1:2] -0.233 -0.251
 $ 1.2: num [1:2] 0.954 -0.266
 $ 1.3: num 1.9
 $ 2.3: num -0.43
 $ 2.4: num [1:2] 1.576 0.162
 $ 2.5: num [1:2] -1.085 0.577
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

37

37

INTRODUCTION TO DATA SCIENCE WITH R



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

16.DEBUGGING

38

38

SOMETHING'S WRONG!

Indications that something's not right

- `message`: A generic notification/diagnostic message produced by the `message` function; execution of the function continues
- `warning`: An indication that something is wrong but not necessarily fatal; execution of the function continues; generated by the `warning` function
- `error`: An indication that a fatal problem has occurred; execution stops; produced by the `stop` function
- `condition`: A generic concept for indicating that something unexpected can occur; programmers can create their own conditions



39

39

SOMETHING'S WRONG!

Warning

```
> log(-1)
[1] NaN
Warning message:
In log(-1) : NaNs produced
```



40

40

SOMETHING'S WRONG!

```
printmessage <- function(x) {
  if(x > 0)
    print("x is greater than zero")
  else
    print("x is less than or equal to zero")
  invisible(x)
}
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

41

41

SOMETHING'S WRONG!

```
> printmessage(1)
[1] "x is greater than zero"
```

```
> printmessage(NA)
Error in if (x > 0) print("x is greater than zero") else print(
  "x is less than or equal to zero") :
  missing value where TRUE/FALSE needed
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

42

42

SOMETHING'S WRONG!

```
printmessage2 <- function(x) {
  if(is.na(x))
    print("x is a missing value!")
  else if(x > 0)
    print("x is greater than zero")
  else
    print("x is less than or equal to zero")
  invisible(x)
}
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

43

43

SOMETHING'S WRONG!

```
> x <- log(-1)
Warning message:
In log(-1) : NaNs produced
```

```
> printmessage2(x)
[1] "x is a missing value!"
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

44

44

SOMETHING'S WRONG!

How do you know that something is wrong with your function?

- What was your input? How did you call the function?
- What were you expecting? Output, messages, other results?
- What did you get?
- How does what you get differ from what you were expecting?
- Were your expectations correct in the first place?
- Can you reproduce the problem (exactly)?



45

45

DEBUGGING TOOLS IN R

The primary tools for debugging functions in R are

- `traceback`: prints out the function call stack after an error occurs; does nothing if there's no error
- `recover`: allows you to modify the error behavior so that you can browse the function call stack
- `debug`: flags a function for "debug" mode which allows you to step through execution of a function one line at a time
- `browser`: suspends the execution of a function wherever it is called and puts the function in debug mode
- `trace`: allows you to insert debugging code into a function at specific places



46

46

traceback

```
> mean(k)
Error in mean(k) : object 'k' not found
> traceback()
1: mean(k)
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

47

47

traceback

```
> lm(y ~ x)
Error in model.frame.default(formula = y ~ x, drop.unused.levels = TRUE) :
  invalid type (list) for variable 'y'
> traceback()
5: model.frame.default(formula = y ~ x, drop.unused.levels = TRUE)
4: stats::model.frame(formula = y ~ x, drop.unused.levels = TRUE)
3: eval(mf, parent.frame())
2: eval(mf, parent.frame())
1: lm(y ~ x)
```



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

48

48

debug

```
> debug(lm)
> lm(y ~ x)
debugging in: lm(y ~ x)
debug: {
  ret.x <- x
  ret.y <- y
  cl <- match.call()
```

```
  if (!qr)
    z$qr <- NULL
  z
}
```

Browse[2]> |



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

49

49

recover

```
> options(error = recover)
> read.csv("nosuchfile")
Error in file(file, "rt") : cannot open the connection
In addition: Warning message:
In file(file, "rt") :
  cannot open file 'nosuchfile': No such file or directory
```

Enter a frame number, or 0 to exit

```
1: read.csv("nosuchfile")
2: read.table(file = file, header = header,
3: file(file, "rt")
```

Selection: |



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

50

50

DEBUGGING

Summary

- There are three main indications of a problem/condition: `message`, `warning`, `error`
 - only an `error` is fatal
- When analyzing a function with a problem, make sure you can reproduce the problem, clearly state your expectations and how the output differs from your expectation
- Interactive debugging tools `traceback`, `debug`, `browser`, `trace`, and `recover` can be used to find problematic code in functions
- Debugging tools are not a substitute for thinking!



ANALYSIS



STRUCTURE



ALGORITHM



PROCESS



PROGRAMMING



SOLVING



KNOWLEDGE

51